

Flight Delay Prediction

Phase 3 — Exploratory Data Analysis (EDA)

Field	Details
Course	COMP4381 – Data Science and Analytics, Spring 2026
Instructor	Ahmed Sabbah

1. Introduction

Flight delays are one of the most serious challenges facing the aviation industry, as they negatively impact the passenger experience and cause financial losses for airlines and airports. This project aims to analyse and predict delay patterns by establishing an integrated data science pipeline.

The dataset used consists of approximately 1.7 million flight records covering nine major airports in the United States during 2024, including information on departure and arrival times, reasons for delays, airlines, and distances.

This report covers the third phase of the project, which is organised as follows: data description, data cleaning and missing value handling, exploratory analysis, and visualisation.

2. Dataset Source

Property	Value
Source	Kaggle – Flight Delay Dataset
Total Records	1,747,627 flights
Total Features	16 columns
Airports Covered	ATL, DFW, JFK, LAX, ORD, BOS, MIA, SEA, SFO
File Path	https://drive.google.com/file/d/1wWvYntzU-XzEKTwzAM891U3jTdzmmy-s/view?usp=drive_link

The flight data set includes scheduled and actual departure and arrival times, airline codes, departure and arrival airports, and the duration of delays in minutes.

3. Import Libraries

We use the standard Python libraries, including:

- **pandas** — data loading, cleaning and manipulation
- **numpy** — numerical operations and array handling
- **matplotlib** / **seaborn** — visualisation

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

pd.set_option('display.max_columns', None)
pd.set_option('display.float_format', lambda x: f'{x:.2f}')

print('[OK] Libraries imported successfully!')
```

[OK] Libraries imported successfully!

4. Load Dataset

The dataset was saved on Google Drive (https://drive.google.com/file/d/1wWvYntzU-XzEKTWzAM891U3jTdzmyy-s/view?usp=drive_link) and loaded in the notebook using its shared file ID. We used **gdown** to download the file as **flight_delays.csv**, then used **pandas.read_csv()** to load it into a DataFrame. This makes the notebook easier to run without manual downloads or path changes.

```
In [2]: import os
try:
    import gdown
except ModuleNotFoundError:
    !pip install -q gdown
    import gdown

file_id = "1wWvYntzU-XzEKTWzAM891U3jTdzmyy-s"
url = f"https://drive.google.com/uc?id={file_id}"
output_path = "flight_delays.csv"

if not os.path.exists(output_path):
    gdown.download(url, output_path, quiet=True)

df = pd.read_csv(output_path)
```

```
print(f"[OK] Data loaded successfully: {df.shape[0]:,} rows, {df.shape[1]:,} columns")
```

[OK] Data loaded successfully: 1,747,627 rows, 16 columns

5. Dataset Structure

We examine the first few rows to understand the data structure before applying any transformations. This enables us to identify obvious formatting issues or unexpected values at an early stage.

```
In [3]: print('=' * 70)
        print('DATASET STRUCTURE – FIRST 5 ROWS')
        print('=' * 70)
        df.head()
```

=====

DATASET STRUCTURE – FIRST 5 ROWS

=====

```
Out[3]:
```

	FlightID	Airline	FlightNumber	Origin	Destination	ScheduledDeparture	ActualDeparture
0	1	United	4558	ORD	MIA	2024-09-01 08:11	2024-09-01 08:11
1	2	Delta	8021	LAX	MIA	2024-09-01 10:25	2024-09-01 10:25
2	3	Southwest	7520	DFW	SFO	2024-09-01 16:53	2024-09-01 16:53
3	4	Delta	2046	ORD	BOS	2024-09-01 14:44	2024-09-01 14:44
4	5	Delta	6049	LAX	SEA	2024-09-01 01:51	2024-09-01 01:51

6. Dataset Dimensions

We confirm the total number of rows and columns to verify the data loaded completely.

```
In [4]: print('=' * 70)
        print('DATASET DIMENSIONS')
        print('=' * 70)
        print(f'Rows      : {df.shape[0]:,}')
        print(f'Columns  : {df.shape[1]:,}')
```

```
=====
DATASET DIMENSIONS
=====
```

```
Rows      : 1,747,627
Columns   : 16
```

7. Column Names & Types

It is necessary to understand the data type of each column before taking any steps to clean the data. Inconsistencies in data types (such as a numeric column stored as an 'object') are a common source of errors.

```
In [5]: print('=' * 70)
        print('COLUMN NAMES')
        print('=' * 70)
        for idx, col in enumerate(df.columns, 1):
            print(f' {idx:2d}. {col}')
```

```
=====
COLUMN NAMES
=====
```

```
1. FlightID
2. Airline
3. FlightNumber
4. Origin
5. Destination
6. ScheduledDeparture
7. ActualDeparture
8. ScheduledArrival
9. ActualArrival
10. DelayMinutes
11. DelayReason
12. Cancelled
13. Diverted
14. AircraftType
15. TailNumber
16. Distance
```

```
In [6]: print('=' * 70)
        print('DATA TYPE DISTRIBUTION')
        print('=' * 70)
        for dtype, count in df.dtypes.value_counts().items():
            print(f' {str(dtype):20s}: {count} column(s)')
```

```
=====
DATA TYPE DISTRIBUTION
=====
```

```
str                : 10 column(s)
int64               : 4 column(s)
bool                : 2 column(s)
```

```
In [7]: print('=' * 70)
        print('COLUMN SUMMARY TABLE')
```

```

print('=' * 70)
desc = pd.DataFrame({
    'Column' : df.columns,
    'Type' : df.dtypes.values,
    'Missing' : df.isnull().sum().values
})
print(desc.to_string(index=False))

```

COLUMN SUMMARY TABLE

```

=====
      Column  Type  Missing
FlightID  int64      0
  Airline   str      0
FlightNumber  int64      0
   Origin   str      0
Destination   str      0
ScheduledDeparture  str      0
ActualDeparture  str      0
ScheduledArrival  str      0
ActualArrival  str      0
DelayMinutes  int64      0
DelayReason   str  468873
Cancelled    bool      0
Diverted     bool      0
AircraftType  str      0
TailNumber   str      0
Distance  int64      0

```

8. Missing Values Analysis

Incomplete data may lead to a skewed model or cause it to fail. In this case, we calculate the total data completeness rate as a percentage and indicate which columns require missing data to be completed or deleted.

A data completeness rate exceeding 95% is generally considered acceptable for tabular datasets.

```

In [8]: completeness = (1 - (df.isnull().sum().sum() / (len(df) * len(df.columns)))) * 100
print(f'  Data Completeness : {completeness:.2f}%')

missing_per_col = df.isnull().sum()
missing_cols = missing_per_col[missing_per_col > 0]
if missing_cols.empty:
    print('  [OK] No missing values detected in any column.')
else:
    print('\n  Columns with missing values:')
    print(missing_cols.to_string())

```

Data Completeness : 98.32%

Columns with missing values:
DelayReason 468873

9. Duplicate Records

Duplicate rows may provide a distortion of the statistical and training summaries. We will inquire again to confirm the presence of duplicate questions and report the number.

```
In [9]: dups = df.duplicated().sum()
print(f' Duplicate rows : {dups:,}')
if dups == 0:
    print(' [OK] No duplicate records found – data is unique.')
else:
    print(f' [!] {dups:,} duplicate rows detected. Consider removing them before m
```

Duplicate rows : 0
[OK] No duplicate records found – data is unique.

10. Descriptive Statistics

We calculate standard summary statistics (count, mean, standard deviation, minimum, quartiles, maximum) for all **numeric columns**. This helps determine the range and spread of delay values and other continuous properties.

```
In [10]: numeric_cols = df.select_dtypes(include=[np.number]).columns
df[numeric_cols].describe()
```

```
Out[10]:
```

	FlightID	FlightNumber	DelayMinutes	Distance
count	1747627.00	1747627.00	1747627.00	1747627.00
mean	873814.00	5001.16	10.00	1549.94
std	504496.60	2885.82	11.83	836.87
min	1.00	1.00	-10.00	100.00
25%	436907.50	2503.00	0.00	825.00
50%	873814.00	5002.00	10.00	1551.00
75%	1310720.50	7499.00	20.00	2274.00
max	1747627.00	9999.00	30.00	3000.00

11. Delay Analysis

The key variable of interest is "delay duration in minutes." Here we will examine the following:

- **Average delay** for all flights
- **Number of delayed flights** (where "delay duration in minutes" > 0)
- **Percentage of delayed flights** to the total number of flights

This provides an initial picture of the frequency and severity of delays.

```
In [11]: delay_stats = df['DelayMinutes'].describe()
delayed   = (df['DelayMinutes'] > 0).sum()
delay_pct = (delayed / len(df)) * 100

print('=' * 70)
print('DELAY SUMMARY')
print('=' * 70)
print(f'  Mean Delay      : {delay_stats["mean"]:.2f} minutes')
print(f'  Median Delay    : {delay_stats["50%"]:.2f} minutes')
print(f'  Max Delay        : {delay_stats["max"]:.2f} minutes')
print(f'  Delayed Flights  : {delayed:,} ({delay_pct:.2f}% of total)')

=====
DELAY SUMMARY
=====

  Mean Delay      : 10.00 minutes
  Median Delay    : 10.00 minutes
  Max Delay       : 30.00 minutes
  Delayed Flights : 1,278,754 (73.17% of total)
```

12. Detecting outliers

We use the **interquartile range (IQR)** to detect outliers in `DelayMinutes`. A value is classified as an outlier if it is less than $Q1 - 1.5 \times IQR$ or greater than $Q3 + 1.5 \times IQR$.

Outlier delay values may result from genuine disruptions (such as severe weather conditions or mechanical failures) or data entry errors

```
In [12]: Q1 = df['DelayMinutes'].quantile(0.25)
Q3 = df['DelayMinutes'].quantile(0.75)
IQR = Q3 - Q1

lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

outliers = ((df['DelayMinutes'] < lower_bound) | (df['DelayMinutes'] > upper_bound))

print('=' * 70)
print('OUTLIER DETECTION - DelayMinutes (IQR Method)')
```

```

print('=' * 70)
print(f' Q1 (25th pct)      : {Q1:.2f} min')
print(f' Q3 (75th pct)      : {Q3:.2f} min')
print(f' IQR                  : {IQR:.2f} min')
print(f' Lower Bound         : {lower_bound:.2f} min')
print(f' Upper Bound         : {upper_bound:.2f} min')
print(f' Outliers Found      : {outliers:,}')

```

```

=====
OUTLIER DETECTION – DelayMinutes (IQR Method)
=====

```

```

Q1 (25th pct)      : 0.00 min
Q3 (75th pct)      : 20.00 min
IQR                 : 20.00 min
Lower Bound        : -30.00 min
Upper Bound        : 50.00 min
Outliers Found     : 0

```

13. Airport Analysis

All historical airports that were either **level airports** or **destination airports** were identified in the dataset.

```

In [13]: origins      = df['Origin'].unique()
destinations = df['Destination'].unique()
all_airports = sorted(set(list(origins) + list(destinations)))

print('=' * 70)
print('AIRPORT COVERAGE')
print('=' * 70)
print(f' Origin airports      : {len(origins)}')
print(f' Destination airports : {len(destinations)}')
print(f' Total unique airports: {len(all_airports)}')
print(f' Airport codes         : {", ".join(all_airports)}')

```

```

=====
AIRPORT COVERAGE
=====

```

```

Origin airports      : 5
Destination airports : 5
Total unique airports: 9
Airport codes        : ATL, BOS, DFW, JFK, LAX, MIA, ORD, SEA, SFO

```

14. Airline Distribution

We examine how flights are distributed among airlines. A highly unbalanced distribution (the dominance of one airline) can lead to bias in the model, so this step is useful in any subsequent decisions regarding stratified sampling.

```
In [14]: airline_counts = df['Airline'].value_counts()

print('=' * 70)
print('AIRLINE DISTRIBUTION')
print('=' * 70)
for airline, count in airline_counts.items():
    pct = (count / len(df)) * 100
    bar = '' * int(pct / 2)
    print(f' {airline:6s} : {count:>9,} ({pct:5.2f}%) {bar}')
```

```
=====
AIRLINE DISTRIBUTION
=====
Southwest : 437,721 (25.05%)
American Airlines : 437,124 (25.01%)
Delta : 436,680 (24.99%)
United : 436,102 (24.95%)
```

15. Data Cleaning

This section focuses on improving the quality of the dataset by identifying and removing duplicate records. Duplicate entries may lead to inaccurate analytical results and affect future machine learning models.

```
In [15]: print("="*70)
print("DATA CLEANING")
print("="*70)

df_clean = df.copy()

print(f"Original Dataset Shape: {df_clean.shape}")

duplicates = df_clean.duplicated().sum()

print(f"\nNumber of Duplicate Records: {duplicates}")

df_clean = df_clean.drop_duplicates()

print(f"Dataset Shape After Removing Duplicates: {df_clean.shape}")
```

```
=====
DATA CLEANING
=====
Original Dataset Shape: (1747627, 16)

Number of Duplicate Records: 0
Dataset Shape After Removing Duplicates: (1747627, 16)
```

16. Missing Values Analysis

Missing values can negatively affect statistical analysis and predictive modeling. This step identifies missing values across all dataset attributes.

```
In [16]: print("="*70)
print("MISSING VALUES ANALYSIS")
print("="*70)

missing_values = df_clean.isnull().sum()

print("Missing Values Per Column:")
print(missing_values[missing_values > 0])
```

```
=====
MISSING VALUES ANALYSIS
=====
Missing Values Per Column:
DelayReason    468873
dtype: int64
```

The output shows the number of missing values in each feature. Features that contain missing values must be handled appropriately before further processing.

17. Missing Values Treatment

Handle missing values to help improve data completeness and maintain the quality of the dataset.

```
In [17]: before_missing = df_clean["DelayReason"].isnull().sum()

df_clean["DelayReason"] = df_clean["DelayReason"].fillna("No Delay/Unknown")

after_missing = df_clean["DelayReason"].isnull().sum()

print(f"Missing values in DelayReason before treatment: {before_missing}")
print(f"Missing values in DelayReason after treatment: {after_missing}")
```

```
Missing values in DelayReason before treatment: 468873
Missing values in DelayReason after treatment: 0
```

Most missing values were in the DelayReason column. We did not remove these rows because they still contain useful flight information. Instead, we replaced the missing values with "No Delay/Unknown" so the dataset stays complete and ready for analysis.

18. Data Quality Verification

Verify that all duplicates and missing values have been successfully handled.

```
In [18]: print("="*70)
print("DATA QUALITY VERIFICATION")
```

```
print("="*70)

print("Remaining Missing Values:")
print(df_clean.isnull().sum().sum())

print("\nRemaining Duplicate Records:")
print(df_clean.duplicated().sum())
```

```
=====
DATA QUALITY VERIFICATION
=====
Remaining Missing Values:
0

Remaining Duplicate Records:
0
```

The dataset was validated following data cleaning to confirm that there are no duplicate records or missing values.

19. Data Preprocessing

Prepare the dataset for feature engineering by standardizing column names and converting date/time columns into appropriate formats.

```
In [19]: print("="*70)
print("DATA PREPROCESSING")
print("="*70)

df_clean.columns = df_clean.columns.str.strip()

date_columns = [
    "ScheduledDeparture",
    "ActualDeparture",
    "ScheduledArrival",
    "ActualArrival"
]

for col in date_columns:
    df_clean[col] = pd.to_datetime(df_clean[col])

print("Date and Time Conversion Completed Successfully")
```

```
=====
DATA PREPROCESSING
=====
Date and Time Conversion Completed Successfully
```

```
In [20]: df_clean["Route"] = df_clean["Origin"] + "-" + df_clean["Destination"]

df_clean["ScheduledHour"] = df_clean["ScheduledDeparture"].dt.hour
df_clean["ArrivalHour"] = df_clean["ScheduledArrival"].dt.hour
df_clean["Weekday"] = df_clean["ScheduledDeparture"].dt.day_name()
```

```
df_clean["Month"] = df_clean["ScheduledDeparture"].dt.month

df_clean[["Route", "ScheduledHour", "ArrivalHour", "Weekday", "Month"]].head()
```

```
Out[20]:
```

	Route	ScheduledHour	ArrivalHour	Weekday	Month
0	ORD-MIA	8	12	Sunday	9
1	LAX-MIA	10	13	Sunday	9
2	DFW-SFO	16	17	Sunday	9
3	ORD-BOS	14	18	Sunday	9
4	LAX-SEA	1	5	Sunday	9

The date and time columns were converted into datetime format, allowing time-based calculations and feature extraction.

20. Feature Engineering: Departure Delay

Calculate the difference between actual and scheduled departure times.

```
In [21]: df_clean["DepartureDelay"] = (
            df_clean["ActualDeparture"] -
            df_clean["ScheduledDeparture"]
        ).dt.total_seconds() / 60

df_clean["DepartureDelay"].describe()
```

```
Out[21]: count    1747627.00
         mean         15.00
         std          8.95
         min          0.00
         25%          7.00
         50%         15.00
         75%         23.00
         max         30.00
         Name: DepartureDelay, dtype: float64
```

DepartureDelay measures the number of minutes a flight departed earlier or later than scheduled.

21. Feature Engineering: Arrival Delay

Calculate the difference between actual and scheduled arrival times.

```
In [22]: df_clean["ArrivalDelay"] = (
            df_clean["ActualArrival"] -
            df_clean["ScheduledArrival"]
        )
```

```
).dt.total_seconds() / 60  
  
df_clean["ArrivalDelay"].describe()
```

```
Out[22]: count    1747627.00  
mean         10.00  
std          11.83  
min          -10.00  
25%           0.00  
50%          10.00  
75%          20.00  
max           30.00  
Name: ArrivalDelay, dtype: float64
```

ArrivalDelay measures the time difference between actual and planned arrival times and provides a useful indicator of operational performance.

22. Feature Engineering: Flight Duration

Calculate the actual flight duration in minutes.

```
In [23]: df_clean["FlightDuration"] = (  
    df_clean["ActualArrival"] -  
    df_clean["ActualDeparture"]  
).dt.total_seconds() / 60  
  
df_clean["FlightDuration"].describe()
```

```
Out[23]: count    1747627.00  
mean         205.01  
std          103.53  
min           20.00  
25%          115.00  
50%          205.00  
75%          295.00  
max           390.00  
Name: FlightDuration, dtype: float64
```

FlightDuration represents the actual travel time of each flight and may help explain delay patterns.

23. Creating Distance Category

Group flights into distance categories for easier analysis.

```
In [24]: def distance_category(distance):  
    if distance < 500:  
        return "Short"  
    elif distance < 1000:  
        return "Medium"
```

```

        else:
            return "Long"

df_clean["DistanceCategory"] = df_clean["Distance"].apply(distance_category)

df_clean["DistanceCategory"].value_counts()

```

```

Out[24]: DistanceCategory
Long      1206064
Medium    300904
Short     240659
Name: count, dtype: int64

```

Flights were categorised into Short, Medium and Long routes based on distance travelled.

24. Creating Delay Category

Classify flights into delay categories based on the number of delay minutes.

```

In [25]: def categorize_delay(delay):
        if delay <= 0:
            return "Early/On Time"
        elif delay <= 15:
            return "Small Delay"
        elif delay <= 30:
            return "Moderate Delay"
        else:
            return "Large Delay"

df_clean["DelayCategory"] = df_clean["DelayMinutes"].apply(categorize_delay)

df_clean["DelayCategory"].value_counts()

```

```

Out[25]: DelayCategory
Small Delay      639437
Moderate Delay   639317
Early/On Time    468873
Name: count, dtype: int64

```

```

In [26]: df_clean["IsDelayed"] = df_clean["DelayMinutes"].apply(lambda x: 1 if x > 15 else 0)

df_clean["IsDelayed"].value_counts()

```

```

Out[26]: IsDelayed
0      1108310
1       639317
Name: count, dtype: int64

```

Flights were grouped into Early/On Time, Minor Delay and Moderate Delay categories to make it easier to analyse delay patterns.

25. Validating the Final Dataset

Review the structure of the final dataset after it has been cleaned and pre-processed.

```
In [27]: print("="*70)
print("FINAL DATASET VALIDATION")
print("="*70)

print("Dataset Shape:")
print(df_clean.shape)

print("\nDataset Information:")
df_clean.info()
```

```
=====
FINAL DATASET VALIDATION
=====
```

Dataset Shape:
(1747627, 27)

Dataset Information:

```
<class 'pandas.DataFrame'>
```

RangeIndex: 1747627 entries, 0 to 1747626

Data columns (total 27 columns):

#	Column	Dtype
0	FlightID	int64
1	Airline	str
2	FlightNumber	int64
3	Origin	str
4	Destination	str
5	ScheduledDeparture	datetime64[us]
6	ActualDeparture	datetime64[us]
7	ScheduledArrival	datetime64[us]
8	ActualArrival	datetime64[us]
9	DelayMinutes	int64
10	DelayReason	str
11	Cancelled	bool
12	Diverted	bool
13	AircraftType	str
14	TailNumber	str
15	Distance	int64
16	Route	str
17	ScheduledHour	int32
18	ArrivalHour	int32
19	Weekday	str
20	Month	int32
21	DepartureDelay	float64
22	ArrivalDelay	float64
23	FlightDuration	float64
24	DistanceCategory	str
25	DelayCategory	str
26	IsDelayed	int64

dtypes: bool(2), datetime64[us](4), float64(3), int32(3), int64(5), str(10)

memory usage: 316.7 MB

The structure of the final dataset was reviewed to ensure that all transformations were successfully applied.

26. Dataset Preview

Display a sample of the final processed dataset.

```
In [28]: df_clean.head()
```

Out[28]:

	FlightID	Airline	FlightNumber	Origin	Destination	ScheduledDeparture	ActualDeparture
0	1	United	4558	ORD	MIA	2024-09-01 08:11:00	2024-09-01 08:11:00
1	2	Delta	8021	LAX	MIA	2024-09-01 10:25:00	2024-09-01 10:25:00
2	3	Southwest	7520	DFW	SFO	2024-09-01 16:53:00	2024-09-01 16:53:00
3	4	Delta	2046	ORD	BOS	2024-09-01 14:44:00	2024-09-01 14:44:00
4	5	Delta	6049	LAX	SEA	2024-09-01 01:51:00	2024-09-01 01:51:00



The preview gives a quick overview of the cleaned and transformed dataset.

27. Saving the Processed Dataset

In [29]:

```
import os

os.makedirs("../data/processed", exist_ok=True)

output_file = "../data/processed/flight_delays_cleaned.csv"

df_clean.to_csv(output_file, index=False)

print("Dataset Saved Successfully")
print(output_file)
```

Dataset Saved Successfully
../data/processed/flight_delays_cleaned.csv

The processed dataset was successfully saved.

28. Summary Statistics

We summarize the main numerical features in the cleaned dataset. The summary statistics help us understand the typical delay time, flight distance, flight duration, and time-related features before creating visualizations.

In [30]:

```
numeric_columns = [
    "DelayMinutes",
    "Distance",
    "FlightDuration",
    "ScheduledHour",
```

```

    "ArrivalHour",
    "IsDelayed"
]

summary_stats = df_clean[numeric_columns].describe()

summary_stats

```

Out[30]:

	DelayMinutes	Distance	FlightDuration	ScheduledHour	ArrivalHour	IsDelayed
count	1747627.00	1747627.00	1747627.00	1747627.00	1747627.00	1747627.00
mean	10.00	1549.94	205.01	11.51	11.50	0.37
std	11.83	836.87	103.53	6.92	6.92	0.48
min	-10.00	100.00	20.00	0.00	0.00	0.00
25%	0.00	825.00	115.00	6.00	5.00	0.00
50%	10.00	1551.00	205.00	12.00	12.00	0.00
75%	20.00	2274.00	295.00	18.00	17.00	1.00
max	30.00	3000.00	390.00	23.00	23.00	1.00

The dataset has 1,747,627 flights. The average delay is 10 minutes, and the median is also 10 minutes, so most delay values are centered around this number. DelayMinutes ranges from -10 to 30 minutes, meaning some flights were early while others had small or moderate delays.

The average distance is about 1549.94, and the average flight duration is about 205 minutes. Also, the mean of IsDelayed is 0.37, which means around 37% of flights are classified as delayed.

In [31]:

```

mean_delay = df_clean["DelayMinutes"].mean()
median_delay = df_clean["DelayMinutes"].median()
std_delay = df_clean["DelayMinutes"].std()
min_delay = df_clean["DelayMinutes"].min()
max_delay = df_clean["DelayMinutes"].max()

print(f"Mean Delay: {mean_delay:.2f} minutes")
print(f"Median Delay: {median_delay:.2f} minutes")
print(f"Standard Deviation: {std_delay:.2f} minutes")
print(f"Minimum Delay: {min_delay:.2f} minutes")
print(f"Maximum Delay: {max_delay:.2f} minutes")

```

```

Mean Delay: 10.00 minutes
Median Delay: 10.00 minutes
Standard Deviation: 11.83 minutes
Minimum Delay: -10.00 minutes
Maximum Delay: 30.00 minutes

```

This dataset shows that the mean and median delay are both 10 minutes. This suggests that delays are generally moderate and not extreme. The minimum delay is -10 minutes and the

maximum is 30 minutes, so this dataset does not include cases of very long delays.

29. Categorical Summary

We summarized the main categorical columns in the cleaned dataset, such as airlines, delay reasons, delay categories, and distance categories. These summaries help us understand how flights are distributed across different groups.

```
In [32]: df_clean["Airline"].value_counts()
```

```
Out[32]: Airline
Southwest          437721
American Airlines  437124
Delta              436680
United            436102
Name: count, dtype: int64
```

The airline distribution is almost balanced. Southwest has the highest number of records with 437,721 flights, followed by American Airlines with 437,124, Delta with 436,680, and United with 436,102. The numbers are very close, so comparing airlines later will be fair because no airline strongly dominates the dataset.

```
In [33]: df_clean["DelayReason"].value_counts()
```

```
Out[33]: DelayReason
No Delay/Unknown    468873
Air Traffic Control  426488
Maintenance         426168
Weather            426098
Name: count, dtype: int64
```

The DelayReason column shows that 468,873 flights have "No Delay/Unknown", which means no clear delay reason was recorded. The recorded delay reasons are also very close in count: Air Traffic Control has 426,488 records, Maintenance has 426,168, and Weather has 426,098. This means the delay reasons are nearly balanced in the dataset.

```
In [34]: df_clean["DelayCategory"].value_counts()
```

```
Out[34]: DelayCategory
Small Delay        639437
Moderate Delay     639317
Early/On Time     468873
Name: count, dtype: int64
```

The DelayCategory output shows that most flights fall into Small Delay and Moderate Delay. Small Delay has 639,437 flights, while Moderate Delay has 639,317 flights. Early/On Time flights are fewer, with 468,873 records. This shows that the dataset contains many delayed flights, but the delays are mostly small or moderate.

```
In [35]: df_clean["DistanceCategory"].value_counts()
```

```
Out[35]: DistanceCategory
Long      1206064
Medium    300904
Short     240659
Name: count, dtype: int64
```

The DistanceCategory output shows that most flights are Long flights, with 1,206,064 records. Medium flights have 300,904 records, and Short flights have 240,659 records. This means the dataset is dominated by long-distance flights, so distance-related results should be interpreted with this imbalance in mind.

30. Visualizations

We used visualizations for better understanding of the patterns in the cleaned dataset. The charts focus on delay minutes, delay categories, airlines, delay reasons, distance categories, and the relationship between distance and delay.


30.1 Distribution of Delay Minutes:

```
In [36]: import os

figures_path = "../figures"
os.makedirs(figures_path, exist_ok=True)
plt.figure(figsize=(8, 5))
plt.hist(df_clean["DelayMinutes"], bins=30)
plt.title("Distribution of Delay Minutes")
plt.xlabel("Delay Minutes")
plt.ylabel("Number of Flights")

plt.savefig("../figures/distribution_delay_minutes.png", bbox_inches="tight")

plt.show()
```



The chart shows that delay values are between -10 and 30 minutes. Some delay values, such as -10 minutes, appear more than 80,000 times, which means many flights arrived or departed early. Other delay values are also spread across the range, showing that the dataset includes early flights, small delays, and moderate delays. There are no extreme delay cases in this dataset.


30.2 Distribution of Delay Categories

```
In [37]: delay_category_counts = df_clean["DelayCategory"].value_counts()

plt.figure(figsize=(8, 5))
plt.bar(delay_category_counts.index, delay_category_counts.values)
plt.title("Distribution of Delay Categories")
plt.xlabel("Delay Category")
plt.ylabel("Number of Flights")
plt.xticks(rotation=20)

plt.savefig("../figures/distribution_delay_categories.png", bbox_inches="tight")

plt.show()
```



This chart shows that Small Delay and Moderate Delay are the most common categories, and both have almost the same number of flights, around 650,000 each. Early/On Time flights are lower, with about 470,000 flights. This means most flights in the dataset had some delay, but the delay was usually small or moderate, not extreme.


30.3 Average Delay by Airline

```
In [38]: airline_delay = df_clean.groupby("Airline")["DelayMinutes"].mean().sort_values()

plt.figure(figsize=(8, 5))
plt.bar(airline_delay.index, airline_delay.values)
plt.title("Average Delay by Airline")
plt.xlabel("Airline")
plt.ylabel("Average Delay Minutes")
plt.xticks(rotation=20)

plt.savefig("../figures/average_delay_by_airline.png", bbox_inches="tight")

plt.show()
```



This chart shows that the average delay is almost the same for all airlines, around 10 minutes. American Airlines, Delta, United, and Southwest have very close average delay values. This means that, in this dataset, no airline clearly has much higher or lower delay than the others.


30.4 Distribution of Delay Reasons

```
In [39]: delay_reason_counts = df_clean["DelayReason"].value_counts()

plt.figure(figsize=(8, 5))
plt.bar(delay_reason_counts.index, delay_reason_counts.values)
plt.title("Distribution of Delay Reasons")
plt.xlabel("Delay Reason")
plt.ylabel("Number of Flights")
plt.xticks(rotation=25)

plt.savefig("../figures/distribution_delay_reasons.png", bbox_inches="tight")

plt.show()
```




This chart shows that "No Delay/Unknown" has the highest number of flights, with about 470,000 records. The other delay reasons are very close to each other: Air Traffic Control, Maintenance, and Weather each have around 420,000 flights. This means the recorded delay reasons are almost balanced in the dataset, while the unknown/no-delay category is slightly higher.

30.5 Average Delay by Distance Category

```
In [40]: distance_delay = df_clean.groupby("DistanceCategory")["DelayMinutes"].mean()

plt.figure(figsize=(8, 5))
plt.bar(distance_delay.index, distance_delay.values)
plt.title("Average Delay by Distance Category")
plt.xlabel("Distance Category")
plt.ylabel("Average Delay Minutes")

plt.savefig("../figures/average_delay_by_distance_category.png", bbox_inches="tight")
plt.show()
```



This chart shows that the average delay is almost the same for all distance categories. Long, medium, and short flights all have an average delay of about 10 minutes. This means that, in this dataset, flight distance does not seem to have a clear effect on the average delay.

30.6 Delay Distribution by Airline


```
In [41]: plt.figure(figsize=(8, 5))

sns.boxplot(data=df_clean, x="Airline", y="DelayMinutes")

plt.title("Delay Minutes Distribution by Airline")
plt.xlabel("Airline")
plt.ylabel("Delay Minutes")
plt.xticks(rotation=20)

plt.savefig("../figures/delay_distribution_by_airline.png", bbox_inches="tight")

plt.show()
```



This boxplot compares the distribution of delay minutes for each airline. The result shows that all airlines have almost the same delay range, from -10 to 30 minutes. The median delay for each airline is around 10 minutes, and the boxes look very similar. This means that delay behavior is almost the same across United, Delta, Southwest, and American Airlines in this dataset.


30.7 Distance vs Delay Minutes

```
In [42]: sample_df = df_clean.sample(5000, random_state=42)

plt.figure(figsize=(8, 5))
plt.scatter(sample_df["Distance"], sample_df["DelayMinutes"], alpha=0.3)
plt.title("Distance vs Delay Minutes")
plt.xlabel("Distance")
plt.ylabel("Delay Minutes")

plt.savefig("../figures/distance_vs_delay_minutes.png", bbox_inches="tight")

plt.show()
```



This scatter plot compares distance with delay minutes using a sample of the dataset. The points are spread across all distances from about 100 to 3000, and delay values still range between -10 and 30 minutes. There is no clear upward or downward pattern, so distance alone does not seem to strongly explain flight delays in this dataset.


30.8 Flight Duration vs Delay Minutes.

```
In [43]: sample_df = df_clean.sample(5000, random_state=42)

plt.figure(figsize=(8, 5))
plt.scatter(sample_df["FlightDuration"], sample_df["DelayMinutes"], alpha=0.3)
plt.title("Flight Duration vs Delay Minutes")
plt.xlabel("Flight Duration (Minutes)")
plt.ylabel("Delay Minutes")

plt.savefig("../figures/flight_duration_vs_delay_minutes.png", bbox_inches="tight")

plt.show()
```



This scatter plot shows flight duration compared with delay minutes. The points appear in separate groups because the dataset has repeated flight duration ranges. Inside each group, delay values still range from -10 to 30 minutes. There is no clear overall upward or downward trend, so flight duration alone does not seem to strongly explain delays in this dataset.

31. Correlation Analysis

WE checked the relationship between the numerical features in the dataset. So correlation helps us understand whether variables such as distance, flight duration, scheduled hour, or arrival hour are related to flight delays.

```
In [44]: correlation_columns = [  
    "DelayMinutes",  
    "Distance",  
    "FlightDuration",  
    "ScheduledHour",  
    "ArrivalHour",  
    "IsDelayed"  
]  
  
correlation_matrix = df_clean[correlation_columns].corr()  
  
correlation_matrix
```

Out[44]:

	DelayMinutes	Distance	FlightDuration	ScheduledHour	ArrivalHour	IsDel
DelayMinutes	1.00	0.00	0.11	-0.00	0.00	
Distance	0.00	1.00	0.00	0.00	0.00	
FlightDuration	0.11	0.00	1.00	-0.00	-0.00	
ScheduledHour	-0.00	0.00	-0.00	1.00	0.28	
ArrivalHour	0.00	0.00	-0.00	0.28	1.00	
IsDelayed	0.83	-0.00	0.10	-0.00	0.00	



The results show a strong positive correlation between the variable `IsDelayed` and the variable `DelayMinutes`, with a value of 0.83. This is to be expected, as `IsDelayed` is derived from the number of minutes of delay. The variable `FlightDuration` shows a weak positive correlation with `DelayMinutes`, with a value of 0.11, meaning that longer flights may experience slightly longer delays, although this relationship is weak. Meanwhile, the correlation coefficients between the variables `Distance`, `ScheduledHour`, `ArrivalHour`, and the variable `DelayMinutes` are close to 0, meaning they do not show a strong direct relationship to flight delays in this dataset.


```
In [45]: plt.figure(figsize=(8, 6))
plt.imshow(correlation_matrix, aspect="auto")
plt.colorbar(label="Correlation")

plt.xticks(range(len(correlation_columns)), correlation_columns, rotation=45)
plt.yticks(range(len(correlation_columns)), correlation_columns)

plt.title("Correlation Matrix of Numerical Features")

plt.savefig("../figures/correlation_matrix.png", bbox_inches="tight")

plt.show()
```



The heatmap confirms the correlations table visually. The strongest relationship is between DelayMinutes and IsDelayed, which appears clearly with a high correlation around 0.83. There is also a small relationship between ScheduledHour and ArrivalHour, around 0.28. Most other feature relationships are very weak and close to 0, especially Distance with DelayMinutes. This means distance, scheduled hour, and arrival hour don't strongly explain flight delays in this dataset.

32. Key Findings and Interpretation

```
In [46]: print("Key Numerical Findings")
print("-" * 40)

print(f"Average delay: {df_clean['DelayMinutes'].mean():.2f} minutes")
print(f"Median delay: {df_clean['DelayMinutes'].median():.2f} minutes")
print(f"Minimum delay: {df_clean['DelayMinutes'].min():.2f} minutes")
print(f"Maximum delay: {df_clean['DelayMinutes'].max():.2f} minutes")
print(f"Delayed flights percentage: {df_clean['IsDelayed'].mean() * 100:.2f}%")

print("\nTop delay category:")
print(df_clean["DelayCategory"].value_counts().head(1))
```

```

print("\nMost common delay reason:")
print(df_clean["DelayReason"].value_counts().head(1))

print("\nStrongest correlations with DelayMinutes:")
print(correlation_matrix["DelayMinutes"].sort_values(ascending=False))

```

Key Numerical Findings

Average delay: 10.00 minutes
 Median delay: 10.00 minutes
 Minimum delay: -10.00 minutes
 Maximum delay: 30.00 minutes
 Delayed flights percentage: 36.58%

Top delay category:
 DelayCategory
 Small Delay 639437
 Name: count, dtype: int64

Most common delay reason:
 DelayReason
 No Delay/Unknown 468873
 Name: count, dtype: int64

Strongest correlations with DelayMinutes:
 DelayMinutes 1.00
 IsDelayed 0.83
 FlightDuration 0.11
 ArrivalHour 0.00
 Distance 0.00
 ScheduledHour -0.00
 Name: DelayMinutes, dtype: float64

The results show that the average delay is 10 minutes, and the median delay is also 10 minutes. This means that delays in the dataset are centered around 10 minutes. The delay values are between -10 and 30 minutes, so the dataset includes early flights and moderate delays, but no cases of extreme delay.

Around 36.58% of flights are classified as delayed. The most common delay category is 'SmallDelay', with 639,437 flights. The most common reason for delay is 'No Delay/Unknown', with 468,873 records.

The correlation shows that the strongest relationship with DelayMinutes is IsDelayed, with a correlation of 0.83. This is to be expected, as IsDelayed was derived from DelayMinutes. FlightDuration has a weak positive correlation with DelayMinutes at 0.11, whilst Distance, ArrivalHour, and Scheduled Hour are close to 0. This means these numerical features do not strongly explain flight delays in this dataset.