

Flight Delay Prediction

Final Project Report — Phase 4

Course: COMP4381 – Data Science and Analytics, Spring 2026

Instructor: Dr. Ahmed Sabbah

Section: 2

Prepared by:

Cileen Antari – 1220944

Shaymaa Dar Mousa – 1222002

Shahd Alhih – 1221817

Abstract

This project focuses on analyzing and predicting flight delays using a large flight dataset. Flight delays are an important problem because they affect passengers, airlines, and airport operations. Delays can cause missed connections, longer waiting times, and extra costs for airlines.

In this project, we used a Kaggle flight delay dataset that contains 1,747,627 flight records and 16 original columns. The dataset covers nine major U.S. airports: ATL, DFW, JFK, LAX, ORD, BOS, MIA, SEA, and SFO. The data includes information such as airline, origin airport, destination airport, scheduled and actual times, delay minutes, delay reason, aircraft type, and distance.

Before building the model, we cleaned and prepared the data. We checked for duplicate records, handled missing values in the `DelayReason` column, converted date and time columns, and created new useful features such as route, weekday, month, scheduled hour, distance category, delay category, and `IsDelayed`.

After preprocessing, we explored the data using summary statistics and visualizations to understand the main delay patterns. Then, we used a Decision Tree classifier to predict whether a flight is delayed or not. We chose this model because it is simple, easy to understand, and suitable for classification problems.

The final goal of this project is to understand flight delay patterns and build a simple prediction model that can classify flights as delayed or not delayed based on the available flight information.

1. Introduction

The aviation industry is an important part of transportation because it connects people, cities, and businesses. Every day, many passengers depend on flights to travel for work, study, tourism, and personal reasons. However, flight delays are still one of the most common problems in air travel.

A flight delay can affect passengers in many ways. It can cause longer waiting times, missed connections, and changes in travel plans. It can also affect airlines because delays may lead to extra costs, schedule problems, and lower customer satisfaction. For this reason, studying flight delays is useful for understanding what factors may be related to delays and how they can be predicted.

In this project, we used data science techniques to analyze flight delay data and build a simple prediction model. The dataset contains flight information such as airline, origin airport, destination airport, scheduled and actual times, delay minutes, delay reasons, aircraft type, and distance. We cleaned the data, created new features, explored the data using visualizations, and then trained a machine learning model to predict whether a flight is delayed or not.

The main goal of this project is to understand delay patterns and use the available flight information to classify flights as delayed or not delayed. This helps show how data analysis and machine learning can be applied to a real-world aviation problem.

2. Background and Domain Overview

Flight delays are a common issue in the aviation industry. A delay happens when a flight leaves or arrives later than its scheduled time. In many aviation studies, a flight is usually considered delayed if it is more than 15 minutes late. Delays can happen for different reasons, such as weather conditions, air traffic control, aircraft maintenance, airport congestion, or delays from previous flights.

The domain of this project is aviation data analytics. In this domain, flight records are studied to understand the factors that may affect flight punctuality. Each row in the dataset represents one flight record. The dataset includes useful information such as the airline, flight number, origin airport, destination airport, scheduled departure and arrival times,

actual departure and arrival times, delay minutes, delay reason, aircraft type, tail number, and distance.

Using this type of data can help us find patterns in flight delays. For example, we can compare delays between airlines, airports, delay reasons, and distance categories. We can also create new features from the date and time columns, such as the weekday, month, and scheduled hour, to better understand when delays are more likely to happen.

Machine learning can also be used in this domain to build a prediction model. In this project, the model predicts whether a flight is delayed or not delayed. This makes the project a classification problem, because the output is a class label instead of a continuous number. The prediction target used in this project is `IsDelayed`, where `1` means the flight is delayed and `0` means the flight is not delayed.

3. Literature Review Summary

Several recent studies have used machine learning to analyze and predict flight delays. These studies show that flight delay prediction is an important problem because delays affect passengers, airlines, and airport operations.

Islam et al. (2025) studied flight delay prediction using flight and weather data. Their study showed that weather features such as temperature, visibility, and wind speed can affect flight delays. They also used explainable AI methods to understand which factors were important in the prediction.

Snell et al. (2026) used U.S. flight data with weather data and tested different machine learning models. Their results showed that ensemble models such as Random Forest performed well, especially when the dataset imbalance was handled using techniques such as SMOTE.

AlBassam and AlShahrani (2025) compared several machine learning algorithms for flight delay prediction. They found that Decision Tree and Random Forest models gave strong results. This study is useful for our project because it supports using a Decision Tree classifier as a simple and understandable model.

Dai (2024) proposed a hybrid machine learning model using aviation big data. The study showed that weather conditions, airport congestion, and historical delay records are important factors in predicting delays. Mamdouh et al. (2023) also highlighted the importance of previous flight delays and schedule-based features in improving prediction performance.

Overall, the reviewed studies show that machine learning is useful for predicting flight delays. They also show that preprocessing, feature engineering, and selecting meaningful

features are important steps before building the model. Based on these studies, our project uses a Decision Tree classifier because it is suitable for classification, easy to interpret, and appropriate for the project level.

4. Dataset Source and Description

The dataset used in this project is a flight delay dataset from Kaggle. It contains 1,747,627 flight records and 16 original columns. Since the dataset file is large, it was uploaded to Google Drive and loaded directly in the notebook using its shared file ID.

Dataset link: https://drive.google.com/file/d/1wWvYntzU-XzEKTWzAM891U3jTdzmyy-s/view?usp=drive_link

The dataset focuses on nine major U.S. airports: ATL, DFW, JFK, LAX, ORD, BOS, MIA, SEA, and SFO. Based on the instructor's clarification, covering at least five major cities or airports was accepted for the project scope. Since this dataset covers nine airports, it satisfies this scope requirement.

Each row in the dataset represents one flight. The original columns include flight ID, airline, flight number, origin airport, destination airport, scheduled departure time, actual departure time, scheduled arrival time, actual arrival time, delay minutes, delay reason, cancellation status, diverted status, aircraft type, tail number, and distance.

This dataset is suitable for the project because it has a large number of records and includes useful flight information that can be used for analysis and prediction. It also allows us to create additional features from time, route, distance, and delay information.

5. Data Loading

In this step, we import the required Python libraries and load the dataset into the notebook. The dataset is downloaded from Google Drive using its file ID, then loaded using `pandas`. This makes the notebook easier to run because the user does not need to download the file manually or change the file path.

```
In [1]: !pip install scikit-learn
```

Requirement already satisfied: scikit-learn in C:\Users\ayman\anaconda3\envs\DS2026\Lib\site-packages (1.9.0)
Requirement already satisfied: numpy>=1.24.1 in C:\Users\ayman\anaconda3\envs\DS2026\Lib\site-packages (from scikit-learn) (2.4.4)
Requirement already satisfied: scipy>=1.10.0 in C:\Users\ayman\anaconda3\envs\DS2026\Lib\site-packages (from scikit-learn) (1.17.1)
Requirement already satisfied: joblib>=1.4.0 in C:\Users\ayman\anaconda3\envs\DS2026\Lib\site-packages (from scikit-learn) (1.5.3)
Requirement already satisfied: narwhals>=2.0.1 in C:\Users\ayman\anaconda3\envs\DS2026\Lib\site-packages (from scikit-learn) (2.22.1)
Requirement already satisfied: threadpoolctl>=3.5.0 in C:\Users\ayman\anaconda3\envs\DS2026\Lib\site-packages (from scikit-learn) (3.6.0)

```
In [2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import os

pd.set_option('display.max_columns', None)
pd.set_option('display.float_format', lambda x: f'{x:.2f}')

print('[OK] Libraries imported successfully!')
```

[OK] Libraries imported successfully!

```
In [3]: try:
import gdown
except ModuleNotFoundError:
!pip install -q gdown
import gdown

file_id = "1wWvYntzU-XzEKTWzAM891U3jTdzmyy-s"
url = f"https://drive.google.com/uc?id={file_id}"
output_path = "flight_delays.csv"

if not os.path.exists(output_path):
    gdown.download(url, output_path, quiet=True)

df = pd.read_csv(output_path)

print(f"[OK] Data loaded successfully: {df.shape[0]:,} rows, {df.shape[1]:,} columns")
```

[OK] Data loaded successfully: 1,747,627 rows, 16 columns

6. Dataset Structure and Initial Understanding

After loading the dataset, we first checked its structure to understand the available columns, data types, missing values, and a small preview of the records. This step is important because

it helps us know what kind of cleaning and preprocessing is needed before analysis and modeling.

```
In [4]: print("=" * 70)
print("DATASET STRUCTURE – FIRST 5 ROWS")
print("=" * 70)

df.head()
```

```
=====
DATASET STRUCTURE – FIRST 5 ROWS
=====
```

```
Out[4]:
```

	FlightID	Airline	FlightNumber	Origin	Destination	ScheduledDeparture	ActualDeparture
0	1	United	4558	ORD	MIA	2024-09-01 08:11	2024-09-01 08:11
1	2	Delta	8021	LAX	MIA	2024-09-01 10:25	2024-09-01 10:25
2	3	Southwest	7520	DFW	SFO	2024-09-01 16:53	2024-09-01 16:53
3	4	Delta	2046	ORD	BOS	2024-09-01 14:44	2024-09-01 14:44
4	5	Delta	6049	LAX	SEA	2024-09-01 01:51	2024-09-01 01:51

```
In [5]: print("=" * 70)
print("DATASET DIMENSIONS")
print("=" * 70)

print(f"Rows      : {df.shape[0]:,}")
print(f"Columns    : {df.shape[1]:,}")
```

```
=====
DATASET DIMENSIONS
=====
```

```
Rows      : 1,747,627
Columns    : 16
```

```
In [6]: print("=" * 70)
print("COLUMN NAMES")
print("=" * 70)

for idx, col in enumerate(df.columns, 1):
    print(f"{idx:2d}. {col}")
```

```
=====
COLUMN NAMES
=====
```

1. FlightID
2. Airline
3. FlightNumber
4. Origin
5. Destination
6. ScheduledDeparture
7. ActualDeparture
8. ScheduledArrival
9. ActualArrival
10. DelayMinutes
11. DelayReason
12. Cancelled
13. Diverted
14. AircraftType
15. TailNumber
16. Distance

```
In [7]: print("=" * 70)
        print("COLUMN SUMMARY TABLE")
        print("=" * 70)

        column_summary = pd.DataFrame({
            "Column": df.columns,
            "Data Type": df.dtypes.values,
            "Missing Values": df.isnull().sum().values
        })

        column_summary
```

```
=====
COLUMN SUMMARY TABLE
=====
```

Out[7]:

	Column	Data Type	Missing Values
0	FlightID	int64	0
1	Airline	str	0
2	FlightNumber	int64	0
3	Origin	str	0
4	Destination	str	0
5	ScheduledDeparture	str	0
6	ActualDeparture	str	0
7	ScheduledArrival	str	0
8	ActualArrival	str	0
9	DelayMinutes	int64	0
10	DelayReason	str	468873
11	Cancelled	bool	0
12	Diverted	bool	0
13	AircraftType	str	0
14	TailNumber	str	0
15	Distance	int64	0

The dataset contains 1,747,627 rows and 16 original columns. The columns include flight information, route information, scheduled and actual times, delay information, and aircraft details. The initial check showed that most columns do not contain missing values, while the `DelayReason` column contains missing values that need to be handled during data cleaning.

7. Data Cleaning

After understanding the basic structure of the dataset, we cleaned the data to make it ready for analysis and modeling. In this step, we checked for duplicate records and missing values.

Duplicate records can affect the analysis because they may repeat the same flight more than once. Missing values can also cause problems in the analysis or machine learning model if they are not handled correctly.

```
In [8]: print("=" * 70)
        print("DATA CLEANING")
        print("=" * 70)
```



```

df_clean = df.copy()

print(f"Original Dataset Shape: {df_clean.shape}")

duplicates = df_clean.duplicated().sum()
print(f"\nNumber of Duplicate Records: {duplicates:,}")

df_clean = df_clean.drop_duplicates()

print(f"Dataset Shape After Removing Duplicates: {df_clean.shape}")

```

```

=====
DATA CLEANING
=====
Original Dataset Shape: (1747627, 16)

```

```

Number of Duplicate Records: 0
Dataset Shape After Removing Duplicates: (1747627, 16)

```

The duplicate check showed that there were no duplicate rows in the dataset. This means that each flight record appears only once, so no records were removed during this step.

```

In [9]: print("=" * 70)
        print("MISSING VALUES ANALYSIS")
        print("=" * 70)

        missing_values = df_clean.isnull().sum()

        print("Missing Values Per Column:")
        print(missing_values[missing_values > 0])

```

```

=====
MISSING VALUES ANALYSIS
=====
Missing Values Per Column:
DelayReason    468873
dtype: int64

```

The missing values analysis showed that the only column with missing values is `DelayReason`. These missing values were not removed because the rows still contain important flight information. Removing them would reduce the size of the dataset and delete useful records.

```

In [10]: before_missing = df_clean["DelayReason"].isnull().sum()

        df_clean["DelayReason"] = df_clean["DelayReason"].fillna("No Delay/Unknown")

        after_missing = df_clean["DelayReason"].isnull().sum()

        print(f"Missing values in DelayReason before treatment: {before_missing:,}")
        print(f"Missing values in DelayReason after treatment: {after_missing:,}")

```

```

Missing values in DelayReason before treatment: 468,873
Missing values in DelayReason after treatment: 0

```

The missing values in the `DelayReason` column were replaced with `No Delay/Unknown` . This value was used because some flights may not have a recorded delay reason, or the reason may be unknown. After this step, the dataset had no remaining missing values.

```
In [11]: print("=" * 70)
print("DATA QUALITY VERIFICATION")
print("=" * 70)

print("Remaining Missing Values:")
print(df_clean.isnull().sum().sum())

print("\nRemaining Duplicate Records:")
print(df_clean.duplicated().sum())
```

```
=====
DATA QUALITY VERIFICATION
=====
Remaining Missing Values:
0

Remaining Duplicate Records:
0
```

After cleaning, the dataset was checked again to make sure there were no remaining missing values or duplicate records. The final check confirmed that the dataset is complete and ready for preprocessing.

8. Data Preprocessing and Feature Engineering

After cleaning the dataset, we prepared it for analysis and modeling. The main preprocessing step was converting the date and time columns from text format into datetime format. This is important because it allows us to extract useful time-based features such as hour, weekday, and month.

Feature engineering was also used to create new meaningful columns from the original data. These new features help us understand the dataset better and make it more useful for the machine learning model.

```
In [12]: print("=" * 70)
print("DATA PREPROCESSING")
print("=" * 70)

df_clean.columns = df_clean.columns.str.strip()

date_columns = [
    "ScheduledDeparture",
    "ActualDeparture",
```

```

    "ScheduledArrival",
    "ActualArrival"
]

for col in date_columns:
    df_clean[col] = pd.to_datetime(df_clean[col])

print("Date and time conversion completed successfully.")

```

```

=====
DATA PREPROCESSING
=====
Date and time conversion completed successfully.

```

The scheduled and actual departure and arrival columns were converted into datetime format. This makes it easier to calculate delays, flight duration, and extract time-related features.

```

In [13]: df_clean["Route"] = df_clean["Origin"] + "-" + df_clean["Destination"]
df_clean["ScheduledHour"] = df_clean["ScheduledDeparture"].dt.hour
df_clean["ArrivalHour"] = df_clean["ScheduledArrival"].dt.hour
df_clean["Weekday"] = df_clean["ScheduledDeparture"].dt.day_name()
df_clean["Month"] = df_clean["ScheduledDeparture"].dt.month

df_clean[["Route", "ScheduledHour", "ArrivalHour", "Weekday", "Month"]].head()

```

```

Out[13]:

```

	Route	ScheduledHour	ArrivalHour	Weekday	Month
0	ORD-MIA	8	12	Sunday	9
1	LAX-MIA	10	13	Sunday	9
2	DFW-SFO	16	17	Sunday	9
3	ORD-BOS	14	18	Sunday	9
4	LAX-SEA	1	5	Sunday	9

Several time and route features were created. The `Route` feature combines the origin and destination airports. The `ScheduledHour`, `ArrivalHour`, `Weekday`, and `Month` features were extracted from the datetime columns. These features can help show whether delay patterns change depending on route or time.

```

In [14]: df_clean["DepartureDelay"] = (
    df_clean["ActualDeparture"] - df_clean["ScheduledDeparture"]
).dt.total_seconds() / 60

df_clean["ArrivalDelay"] = (
    df_clean["ActualArrival"] - df_clean["ScheduledArrival"]
).dt.total_seconds() / 60

df_clean["FlightDuration"] = (
    df_clean["ActualArrival"] - df_clean["ActualDeparture"]
).dt.total_seconds() / 60

```

```
df_clean[["DepartureDelay", "ArrivalDelay", "FlightDuration"]].describe()
```

Out[14]:

	DepartureDelay	ArrivalDelay	FlightDuration
count	1747627.00	1747627.00	1747627.00
mean	15.00	10.00	205.01
std	8.95	11.83	103.53
min	0.00	-10.00	20.00
25%	7.00	0.00	115.00
50%	15.00	10.00	205.00
75%	23.00	20.00	295.00
max	30.00	30.00	390.00

The `DepartureDelay` and `ArrivalDelay` features measure the difference between the actual and scheduled times in minutes. The `FlightDuration` feature measures the actual duration of each flight. These features are useful for analysis because they describe how flight timing differs from the schedule.

```
In [15]: def distance_category(distance):
          if distance < 500:
              return "Short"
          elif distance < 1000:
              return "Medium"
          else:
              return "Long"

          df_clean["DistanceCategory"] = df_clean["Distance"].apply(distance_category)

          df_clean["DistanceCategory"].value_counts()
```

```
Out[15]: DistanceCategory
Long      1206064
Medium    300904
Short      240659
Name: count, dtype: int64
```

The `DistanceCategory` feature groups flights into Short, Medium, and Long routes. This makes it easier to compare delay patterns based on flight distance.

```
In [16]: def categorize_delay(delay):
          if delay <= 0:
              return "Early/On Time"
          elif delay <= 15:
              return "Small Delay"
          elif delay <= 30:
              return "Moderate Delay"
          else:
```

```

        return "Large Delay"

df_clean["DelayCategory"] = df_clean["DelayMinutes"].apply(categorize_delay)

df_clean["IsDelayed"] = df_clean["DelayMinutes"].apply(lambda x: 1 if x > 15 else 0)

print("Delay Category Counts:")
print(df_clean["DelayCategory"].value_counts())

print("\nTarget Variable Counts:")
print(df_clean["IsDelayed"].value_counts())

```

```

Delay Category Counts:
DelayCategory
Small Delay      639437
Moderate Delay   639317
Early/On Time    468873
Name: count, dtype: int64

```

```

Target Variable Counts:
IsDelayed
0      1108310
1       639317
Name: count, dtype: int64

```

The `DelayCategory` feature was created to group flights based on delay length. The target variable for the machine learning model is `IsDelayed`. A flight is labeled as delayed (`1`) if its delay is more than 15 minutes, and not delayed (`0`) otherwise. This changes the problem into a binary classification task.

```

In [17]: print("=" * 70)
          print("FINAL DATASET VALIDATION")
          print("=" * 70)

          print("Dataset Shape:")
          print(df_clean.shape)

          print("\nDataset Information:")
          df_clean.info()

```

```
=====
FINAL DATASET VALIDATION
=====
Dataset Shape:
(1747627, 27)

Dataset Information:
<class 'pandas.DataFrame'>
RangeIndex: 1747627 entries, 0 to 1747626
Data columns (total 27 columns):
#   Column                Dtype
---  -
0   FlightID              int64
1   Airline               str
2   FlightNumber          int64
3   Origin                str
4   Destination           str
5   ScheduledDeparture    datetime64[us]
6   ActualDeparture       datetime64[us]
7   ScheduledArrival      datetime64[us]
8   ActualArrival         datetime64[us]
9   DelayMinutes          int64
10  DelayReason            str
11  Cancelled              bool
12  Diverted               bool
13  AircraftType           str
14  TailNumber             str
15  Distance               int64
16  Route                  str
17  ScheduledHour          int32
18  ArrivalHour            int32
19  Weekday                str
20  Month                  int32
21  DepartureDelay         float64
22  ArrivalDelay           float64
23  FlightDuration         float64
24  DistanceCategory       str
25  DelayCategory          str
26  IsDelayed              int64
dtypes: bool(2), datetime64[us](4), float64(3), int32(3), int64(5), str(10)
memory usage: 316.7 MB
```

After preprocessing and feature engineering, the final dataset contains 27 columns. This means the dataset now has more than 15 meaningful features, which satisfies the project requirement. The dataset is now ready for exploratory analysis and machine learning.

9. Exploratory Data Analysis and Visualizations

In this section, we explored the cleaned dataset using summary statistics and visualizations. The purpose of this step is to understand the main patterns in the data before building the

machine learning model.

We focused on delay minutes, delay categories, airlines, delay reasons, distance categories, and the relationship between distance and delay. These analyses help us understand how delays are distributed and whether some features show clear patterns related to flight delays.

```
In [18]: numeric_columns = [
    "DelayMinutes",
    "Distance",
    "FlightDuration",
    "ScheduledHour",
    "ArrivalHour",
    "IsDelayed"
]

summary_stats = df_clean[numeric_columns].describe()
summary_stats
```

```
Out[18]:
```

	DelayMinutes	Distance	FlightDuration	ScheduledHour	ArrivalHour	IsDelayed
count	1747627.00	1747627.00	1747627.00	1747627.00	1747627.00	1747627.00
mean	10.00	1549.94	205.01	11.51	11.50	0.37
std	11.83	836.87	103.53	6.92	6.92	0.48
min	-10.00	100.00	20.00	0.00	0.00	0.00
25%	0.00	825.00	115.00	6.00	5.00	0.00
50%	10.00	1551.00	205.00	12.00	12.00	0.00
75%	20.00	2274.00	295.00	18.00	17.00	1.00
max	30.00	3000.00	390.00	23.00	23.00	1.00

The summary statistics show that the average delay is about 10 minutes, and the median delay is also 10 minutes. Delay minutes range from -10 to 30 minutes, which means the dataset includes early flights, on-time flights, small delays, and moderate delays, but it does not include very long delay cases.

The average flight distance is about 1,550 miles, and the average flight duration is about 205 minutes. The mean value of `IsDelayed` is around 0.37, which means that about 37% of flights are classified as delayed based on the selected threshold.

```
In [19]: print("Airline Distribution:")
print(df_clean["Airline"].value_counts())

print("\nDelay Reason Distribution:")
print(df_clean["DelayReason"].value_counts())

print("\nDelay Category Distribution:")
```

```
print(df_clean["DelayCategory"].value_counts())

print("\nDistance Category Distribution:")
print(df_clean["DistanceCategory"].value_counts())
```

Airline Distribution:

```
Airline
Southwest      437721
American Airlines  437124
Delta          436680
United         436102
Name: count, dtype: int64
```

Delay Reason Distribution:

```
DelayReason
No Delay/Unknown    468873
Air Traffic Control  426488
Maintenance         426168
Weather            426098
Name: count, dtype: int64
```

Delay Category Distribution:

```
DelayCategory
Small Delay      639437
Moderate Delay   639317
Early/On Time    468873
Name: count, dtype: int64
```

Distance Category Distribution:

```
DistanceCategory
Long      1206064
Medium    300904
Short     240659
Name: count, dtype: int64
```

The categorical summaries show that the airline distribution is almost balanced between Southwest, American Airlines, Delta, and United. This is useful because no airline strongly dominates the dataset.

For delay reasons, the largest group is `No Delay/Unknown`, followed by Air Traffic Control, Maintenance, and Weather. The delay category summary shows that most flights are either small delay or moderate delay, while early or on-time flights are fewer. The distance category summary shows that most flights are long-distance flights, so distance-related results should be interpreted with this in mind.


```
In [20]: import os

figures_path = "../figures"
os.makedirs(figures_path, exist_ok=True)

plt.figure(figsize=(8, 5))
plt.hist(df_clean["DelayMinutes"], bins=30)
plt.title("Distribution of Delay Minutes")
plt.xlabel("Delay Minutes")
```




```
plt.ylabel("Number of Flights")
plt.savefig("../figures/distribution_delay_minutes.png", bbox_inches="tight")
plt.show()
```



The distribution of delay minutes shows that delay values are between -10 and 30 minutes. This means the dataset does not include extreme delays. The chart also shows that flights are spread across early arrival, small delay, and moderate delay values.

```
In [21]: delay_category_counts = df_clean["DelayCategory"].value_counts()


plt.figure(figsize=(8, 5))
plt.bar(delay_category_counts.index, delay_category_counts.values)
plt.title("Distribution of Delay Categories")
plt.xlabel("Delay Category")
plt.ylabel("Number of Flights")
plt.xticks(rotation=20)
plt.savefig("../figures/distribution_delay_categories.png", bbox_inches="tight")
plt.show()
```



The delay category chart shows that Small Delay and Moderate Delay are the most common categories. Early or on-time flights are lower in number. This means that many flights in the dataset experienced some delay, but the delays were usually not extreme.

```
In [22]: airline_delay = df_clean.groupby("Airline")["DelayMinutes"].mean().sort_values()


plt.figure(figsize=(8, 5))
plt.bar(airline_delay.index, airline_delay.values)
plt.title("Average Delay by Airline")
plt.xlabel("Airline")
plt.ylabel("Average Delay Minutes")
plt.xticks(rotation=20)
plt.savefig("../figures/average_delay_by_airline.png", bbox_inches="tight")
plt.show()
```



The average delay by airline chart shows that the average delay is very close across all airlines. American Airlines, Delta, United, and Southwest all have an average delay of around 10 minutes. This means that, in this dataset, no airline clearly has a much higher or much lower average delay than the others.

```
In [23]: delay_reason_counts = df_clean["DelayReason"].value_counts()


plt.figure(figsize=(8, 5))
plt.bar(delay_reason_counts.index, delay_reason_counts.values)
plt.title("Distribution of Delay Reasons")
plt.xlabel("Delay Reason")
plt.ylabel("Number of Flights")
plt.xticks(rotation=25)
plt.savefig("../figures/distribution_delay_reasons.png", bbox_inches="tight")
plt.show()
```



The delay reason chart shows that `No Delay/Unknown` has the highest number of records. The other delay reasons, including Air Traffic Control, Maintenance, and Weather, have very close counts. This shows that the recorded delay reasons are almost balanced in the dataset, but many flights still do not have a clear recorded reason.

```
In [24]: distance_category_counts = df_clean["DistanceCategory"].value_counts()


plt.figure(figsize=(8, 5))
plt.bar(distance_category_counts.index, distance_category_counts.values)
plt.title("Distribution of Distance Categories")
plt.xlabel("Distance Category")
plt.ylabel("Number of Flights")
plt.savefig("../figures/distribution_distance_categories.png", bbox_inches="tight")
plt.show()
```



The distance category chart shows that most flights in the dataset are long-distance flights. Medium and short flights appear less often. This means that the dataset is not fully balanced by distance category, so any distance-related result should be interpreted carefully.

```
In [25]: distance_delay = df_clean.groupby("DistanceCategory")["DelayMinutes"].mean().sort_v


plt.figure(figsize=(8, 5))
plt.bar(distance_delay.index, distance_delay.values)
plt.title("Average Delay by Distance Category")
plt.xlabel("Distance Category")
plt.ylabel("Average Delay Minutes")
plt.savefig("../figures/average_delay_by_distance_category.png", bbox_inches="tight")
plt.show()
```



The average delay by distance category chart compares the delay values for short, medium, and long flights. The results show whether flight distance has a clear relationship with delay time. In this dataset, the average delay values are generally close, so distance category does not appear to strongly change the average delay.

```
In [26]: hour_delay = df_clean.groupby("ScheduledHour")["DelayMinutes"].mean()

plt.figure(figsize=(9, 5))
plt.plot(hour_delay.index, hour_delay.values, marker="o")
plt.title("Average Delay by Scheduled Departure Hour")
plt.xlabel("Scheduled Departure Hour")
plt.ylabel("Average Delay Minutes")
plt.xticks(range(0, 24))
plt.savefig("../figures/average_delay_by_scheduled_hour.png", bbox_inches="tight")
plt.show()
```



The average delay by scheduled hour chart shows how delay changes during different hours of the day. This helps us check if flights scheduled at certain times have higher average delays. The values are close across the hours, which means there is no very strong delay pattern based only on scheduled departure hour in this dataset.

10. Machine Learning Model Preparation

In this section, we prepare the dataset for the machine learning model. The goal is to predict whether a flight is delayed or not using the cleaned and processed dataset.

10.1 Target Variable Selection

The target variable used in this project is `IsDelayed`. This variable was created during preprocessing. A flight is labeled as delayed (`1`) if the delay is more than 15 minutes, and not delayed (`0`) otherwise.

This makes the problem a binary classification problem because the model predicts one of two possible classes: delayed or not delayed.

10.2 Feature Selection and Data Leakage Prevention

Before training the model, we selected the input features carefully. We did not use `DelayMinutes`, `ArrivalDelay`, or `DepartureDelay` as input features because they are

directly related to the target variable. Using these columns could cause data leakage, meaning the model may learn from information that already reveals the answer.

The selected features describe the airline, route, schedule time, distance, and delay reason. These features are more suitable because they help the model learn useful patterns without directly giving it the target answer.

```
In [27]: print("=" * 70)
print("TARGET AND FEATURE SELECTION")
print("=" * 70)

model_df = df_clean.copy()

target = "IsDelayed"

feature_columns = [
    "Airline",
    "Origin",
    "Destination",
    "Distance",
    "ScheduledHour",
    "ArrivalHour",
    "Weekday",
    "Month",
    "DistanceCategory",
    "DelayReason"
]

X = model_df[feature_columns]
y = model_df[target]

print("Target Variable:", target)

print("\nSelected Features:")
for col in feature_columns:
    print("-", col)

print("\nFeature Data Shape:", X.shape)
print("Target Data Shape:", y.shape)

print("\nTarget Distribution:")
print(y.value_counts())

print("\nTarget Distribution Percentage:")
print((y.value_counts(normalize=True) * 100).round(2))
```



```
=====
TARGET AND FEATURE SELECTION
=====
```

Target Variable: IsDelayed

Selected Features:

- Airline
- Origin
- Destination
- Distance
- ScheduledHour
- ArrivalHour
- Weekday
- Month
- DistanceCategory
- DelayReason

Feature Data Shape: (1747627, 10)

Target Data Shape: (1747627,)

Target Distribution:

IsDelayed

0 1108310

1 639317

Name: count, dtype: int64

Target Distribution Percentage:

IsDelayed

0 63.42

1 36.58

Name: proportion, dtype: float64

The target distribution shows the number of delayed and not delayed flights. Checking this distribution is important because it helps us understand whether the classes are balanced or not before training the model.

10.3 Encoding Categorical Features

The selected features include categorical columns such as airline, origin airport, destination airport, weekday, distance category, and delay reason. Since machine learning models need numerical input, these categorical features were converted into numeric columns using one-hot encoding.

```
In [28]: X_encoded = pd.get_dummies(X, drop_first=True)

print("=" * 70)
print("ENCODING COMPLETED")
print("=" * 70)

print("Original Feature Shape:", X.shape)
print("Encoded Feature Shape:", X_encoded.shape)
```

```
X_encoded.head()
```

```
=====
ENCODING COMPLETED
=====
```

```
Original Feature Shape: (1747627, 10)
```

```
Encoded Feature Shape: (1747627, 20)
```

Out[28]:

	Distance	ScheduledHour	ArrivalHour	Month	Airline_Delta	Airline_Southwest	Airline_
0	1031	8	12	9	False	False	
1	1006	10	13	9	True	False	
2	2980	16	17	9	False	True	
3	1408	14	18	9	True	False	
4	2298	1	5	9	True	False	



10.4 Train-Test Split

The dataset was split into training and testing sets. We used 80% of the data for training and 20% for testing. The training set is used to train the model, while the testing set is used to evaluate how well the model performs on unseen data.

The `stratify=y` option was used to keep the same class distribution in both the training and testing sets. This makes the evaluation more reliable.

```
In [29]: from sklearn.model_selection import train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(
    X_encoded,
    y,
    test_size=0.20,
    random_state=42,
    stratify=y
)
```

```
print("=" * 70)
print("TRAIN / TEST SPLIT")
print("=" * 70)
```

```
print("X_train Shape:", X_train.shape)
print("X_test Shape :", X_test.shape)
print("y_train Shape:", y_train.shape)
print("y_test Shape :", y_test.shape)
```

```
print("\nTraining Target Distribution:")
print(y_train.value_counts(normalize=True).round(3))
```

```
print("\nTesting Target Distribution:")
print(y_test.value_counts(normalize=True).round(3))
```

```
=====
TRAIN / TEST SPLIT
=====
```

```
X_train Shape: (1398101, 20)
X_test Shape : (349526, 20)
y_train Shape: (1398101,)
y_test Shape : (349526,)
```

```
Training Target Distribution:
IsDelayed
0    0.63
1    0.37
Name: proportion, dtype: float64
```

```
Testing Target Distribution:
IsDelayed
0    0.63
1    0.37
Name: proportion, dtype: float64
```

At this point, the dataset is ready for model training. The target variable, selected features, encoded feature matrix, and train-test split have all been prepared. The next step is to train the Decision Tree classifier and evaluate its performance using the testing set.

11. Model Evaluation Results

12. Discussion

13. Limitations

14. Conclusion

15. References

In []: